

[nagels]

Graph Meetup Mumbai

# Roll your own Graph-driven MCP Server

# About Us: Ghlen Nagels



## **CEO of Nagels IT Consulting**

Digital Nomad, loves to run, enjoys craft beers

**Email:** ghlen@nagels.tech

**Website:** <https://nagels.tech>



# About Us: Nabeel Parkar

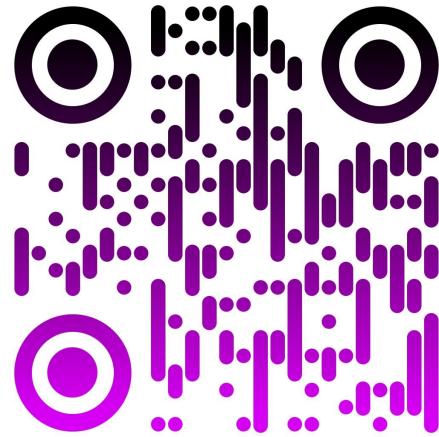


**Open Source Engineer**

South African somehow ended up in India

**Email:** nabeel@ex3.dev

**Website:** <https://ex3.dev>



# About Nagels IT Consulting

- Human First Developer Teams
- Focus on International SMEs
- Specializes in Cloud Applications and Graph Technologies
- Consulting, Training & Development
- [ We're Hiring ]



# The Game Plan

- **Neo4J:** birds eye view
- **MCP:** The high level
- **Developing:** Driver concepts
- **Showoff time:** Combining it all!

# Neo4J: Birds Eye View

# Mumbai Metropolitan Region Rail Map



Copyright © 2025 Nagels

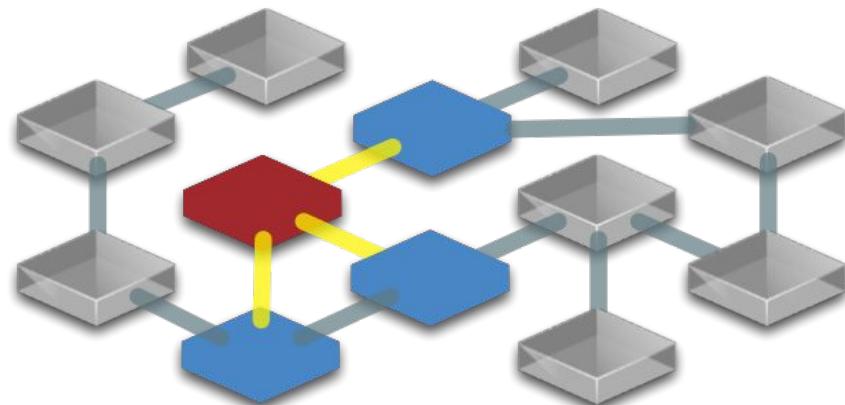
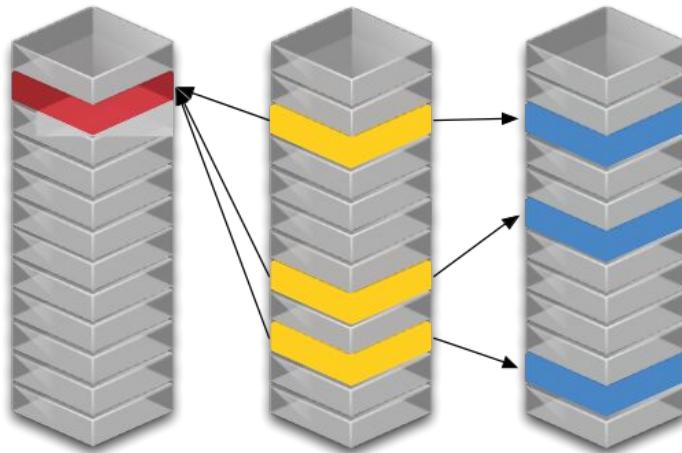
[nagels]

[n]

# Relational Database

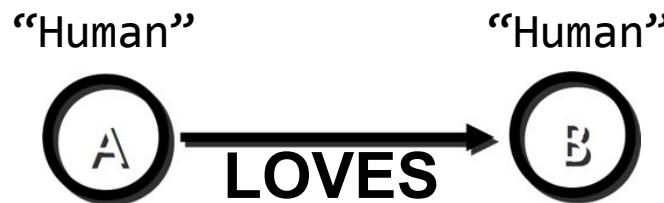
# Graph Database

Let's Build Your Tech Together



[n]

# Cypher Query Language example



## Cypher

```
(A : Human) - [:LOVES] -> (B : Human)
```

# Cypher Basic Query Structure

## Basic query structure

(OPTIONAL) MATCH graph pattern

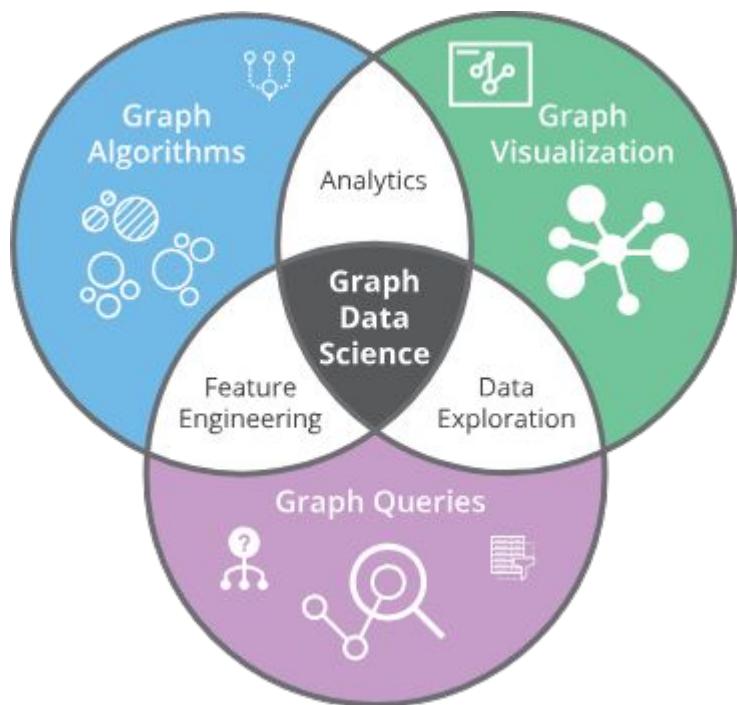
WHERE predicates

RETURN expression

# What to know about Graphs

- Highly **efficient** with connected data
- Extremely friendly to **data science**
- **Different paradigm** from relational DB
- **Cypher** is the common query language
- **Easy to understand**, hard to master
- One of the **fastest growing** technologies

# Learn Graphs at the Graph Academy



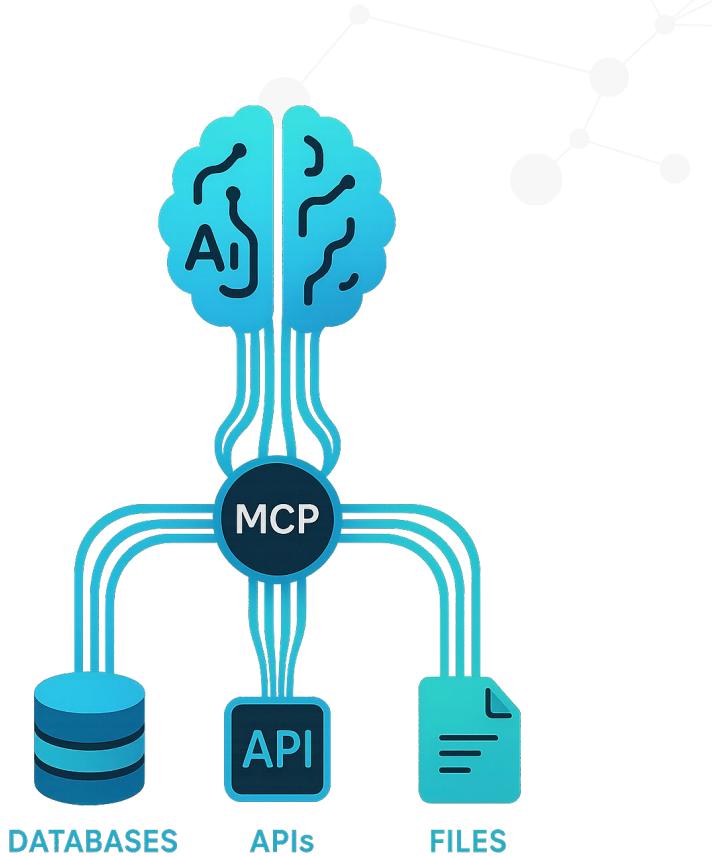
<https://graphacademy.neo4j.com/>



# MCP: The High Level

# Model Context Protocol (MCP)

Connecting AI to Real-World Data

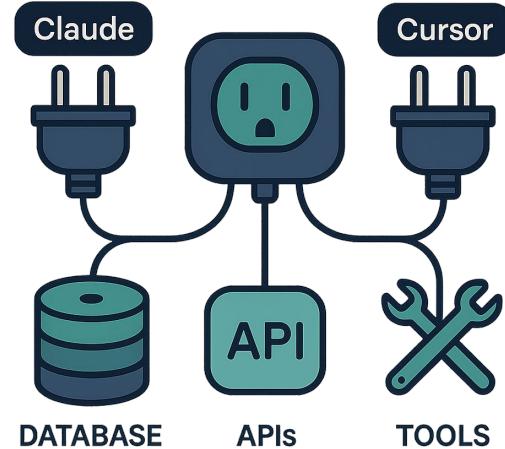


# What is MCP:

The Model Context Protocol (MCP) is an open standard that enables AI assistants to securely connect to external data sources and tools in real-time.

Think of it as...

A **universal adapter** that lets AI assistants “plug into” various data sources and tools without requiring custom integrations for each AI platform.



# Core Benefits of MCP

## Standardization

Write once, use with many AI platforms (Claude, Cursor, VS Code, etc.)

## Security

User-controlled access with permission prompts

## Real-time Data

Access to live, current information (no stale training data)

## Extensibility

Easy to add new tools and data sources

# Industry Adoption

## Major AI Players

- Google
- Microsoft
- AWS
- OpenAI

## Agent Frameworks

- LangChain
- Pydantic-AI
- CrewAI
- LlamalIndex

## Developer Tools

- Claude Desktop
- Cursor AI IDE
- VsCode

## Growing Ecosystem

**Thousands** of community-created MCP servers

# The Problem MCP Solves

## Data Silos

AI models can't access your live business data

## Custom Integrations

Each AI platform requires separate tool development

## Security Concerns

No standardized way to control AI access to sensitive data

## Stale Information

AI responses based on outdated training data

# Real-World Use Cases

## Customer Support

- Query Neo4j knowledge graphs for troubleshooting
- Real-time customer history access
- Dynamic problem resolution

## Development & DevOps

- Database schema queries and diagnostics
- Automated infrastructure management
- Architecture-based code generation

## Business Intelligence

- Natural language queries against live graph analytics
- Path analysis between business entities
- Customer churn prediction

## Enterprise Knowledge Management

- Live documentation and policy access
- Real-time compliance checking
- Dynamic content generation

# MCP Components: Tools & Resources

## Tools - Functions that AI can execute

- Query Neo4j knowledge graphs for troubleshooting
- Real-time customer history access
- Dynamic problem resolution

## Resources - Static or semi-static data for context

- Natural language queries against live graph analytics
- Path analysis between business entities
- Customer churn prediction

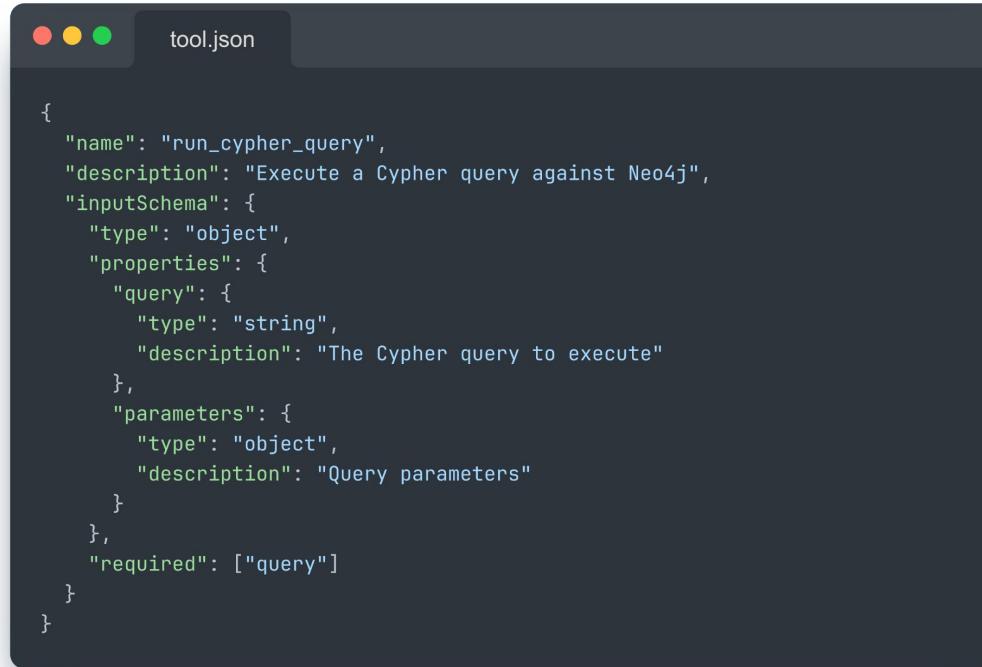
## Simple Distinction

**Tools:** "What can I do?" (Actions)

**Resources:** "What do I need to know?" (Context)



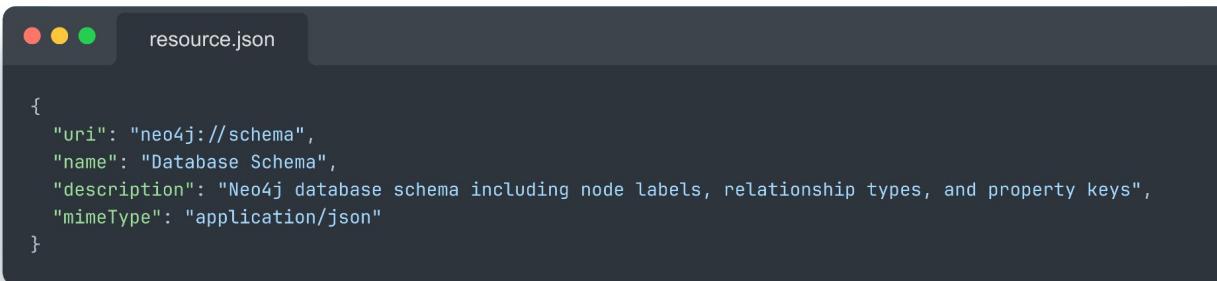
# Tools Example: Neo4j Cypher Query



A screenshot of a Mac OS X application window titled "tool.json". The window contains the following JSON code:

```
{  
  "name": "run_cypher_query",  
  "description": "Execute a Cypher query against Neo4j",  
  "inputSchema": {  
    "type": "object",  
    "properties": {  
      "query": {  
        "type": "string",  
        "description": "The Cypher query to execute"  
      },  
      "parameters": {  
        "type": "object",  
        "description": "Query parameters"  
      }  
    },  
    "required": ["query"]  
  }  
}
```

# Resources Example: Database Schema



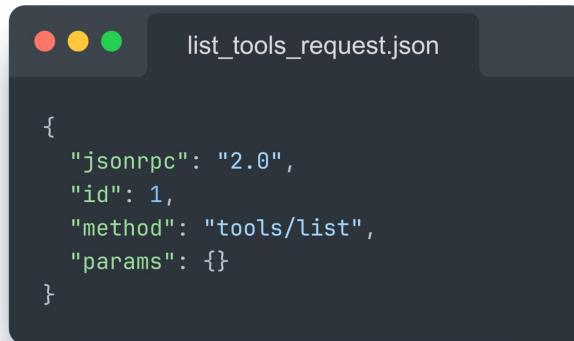
A screenshot of a Mac OS X-style file browser window titled "resource.json". The window has the standard red, yellow, and green close buttons at the top left. The main content area shows a JSON object:

```
{  
  "uri": "neo4j://schema",  
  "name": "Database Schema",  
  "description": "Neo4j database schema including node labels, relationship types, and property keys",  
  "mimeType": "application/json"  
}
```

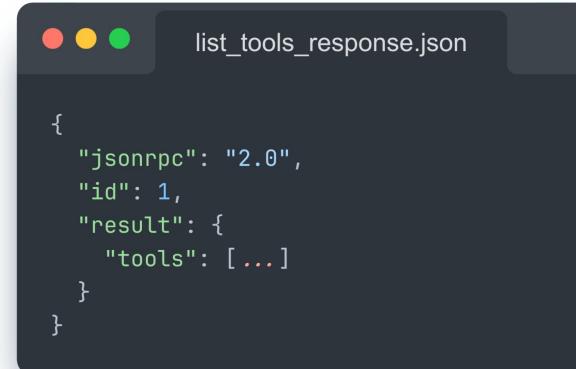
# How MCP Works: JSON-RPC 2.0

## Foundation Protocol

- UTF-8 encoded messages
- Request/response pattern with unique IDs
- Support for notifications (no response expected)
- Batch operations for multiple requests



```
{  
    "jsonrpc": "2.0",  
    "id": 1,  
    "method": "tools/list",  
    "params": {}  
}
```



```
{  
    "jsonrpc": "2.0",  
    "id": 1,  
    "result": {  
        "tools": [ ... ]  
    }  
}
```

# Transport Mechanisms

## STDIO (Standard Input/Output)

- Local development & desktop apps
- Client launches server as subprocess
- Messages via stdin/stdout

## HTTP+SSE (Legacy - 2024)

- Deprecated but widely supported
- SSE for server-to-client messages
- HTTP POST for client-to-server

## Streamable HTTP (New - 2025)

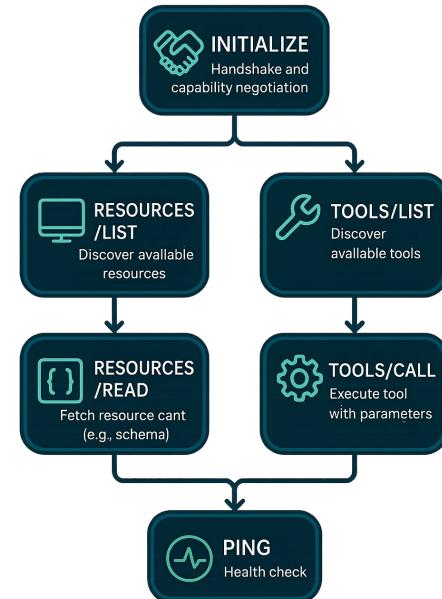
- Single GET and POST HTTP endpoint
- Optional SSE streaming
- Session management with headers

# Message Types in MCP

## Core Operations

1. **Initialize**
  - Handshake and capability negotiation
2. **Resources/List**
  - Discover available resources
3. **Resources/Read**
  - Fetch resource content (e.g., schema)
4. **Tools/List**
  - Discover available tools
5. **Tools/Call**
  - Execute tool with parameters
6. **Ping**
  - Health check

## MCP OPERATIONS



## OAuth 2.1 Integration

- **Standard Protocol:** Based on OAuth 2.1 with PKCE for security
- **Transport Level:** Authorization at HTTP transport layer
- **Optional Implementation:** Not required for all MCP servers

## Security Features

- **HTTPS Required:** All authorization endpoints must use HTTPS
- **Token Validation:** Access tokens validated on every request
- **PKCE Required:** Proof Key for Code Exchange for all clients
- **Secure Storage:** OAuth 2.1 best practices for token handling

## Grant Types Supported

- **Authorization Code:** User-based authentication (human end users)
- **Client Credentials:** Application-to-application authentication
- **Dynamic Client Registration:** Automatic client registration (RFC7591)

## Discovery & Fallbacks

- **Metadata Discovery:** OAuth 2.0 Authorization Server Metadata (RFC8414)
- **Default Endpoints:** `/authorize`, `/token`, `/register` fallbacks
- **Third-Party Auth:** Support for delegated authorization servers

## Validation & Registry

- **Reference Client Implementations:** High-quality AI applications demonstrating protocol features
- **Compliance Test Suites:** Automated verification for clients, servers, and SDKs
- **MCP Registry:** Centralized discovery and API layer for third-party marketplaces
- **Developer Ecosystem:** Tools for confident MCP implementation

## Agent Enhancements

- **Agent Graphs:** Complex agent topologies through namespacing
- **Interactive Workflows:** Human-in-the-loop experiences with granular permissions
- **Direct Communication:** Standardized patterns for end-user interaction

## Multimodality & Streaming

- **Additional Media Types:** Video and other modalities beyond text
- **Streaming Support:** Multipart, chunked messages for interactive experiences
- **Bidirectional Communication:** Real-time, interactive data flows

## Governance & Community

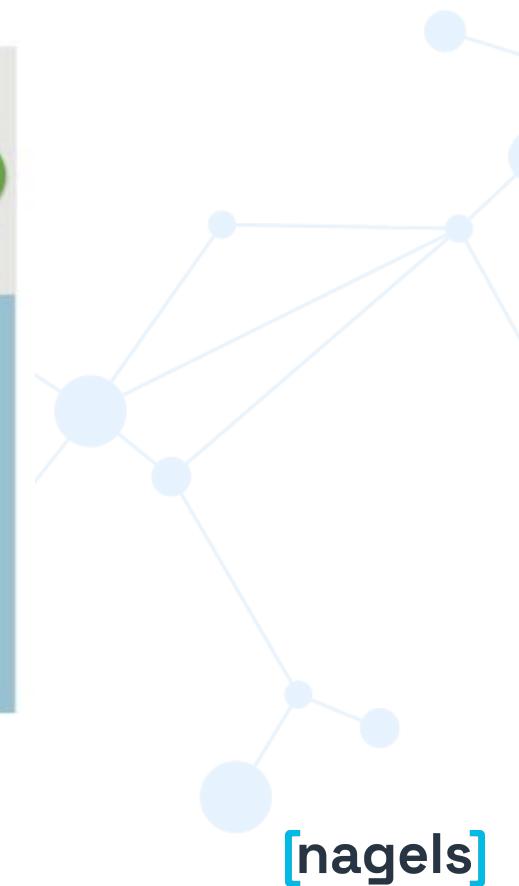
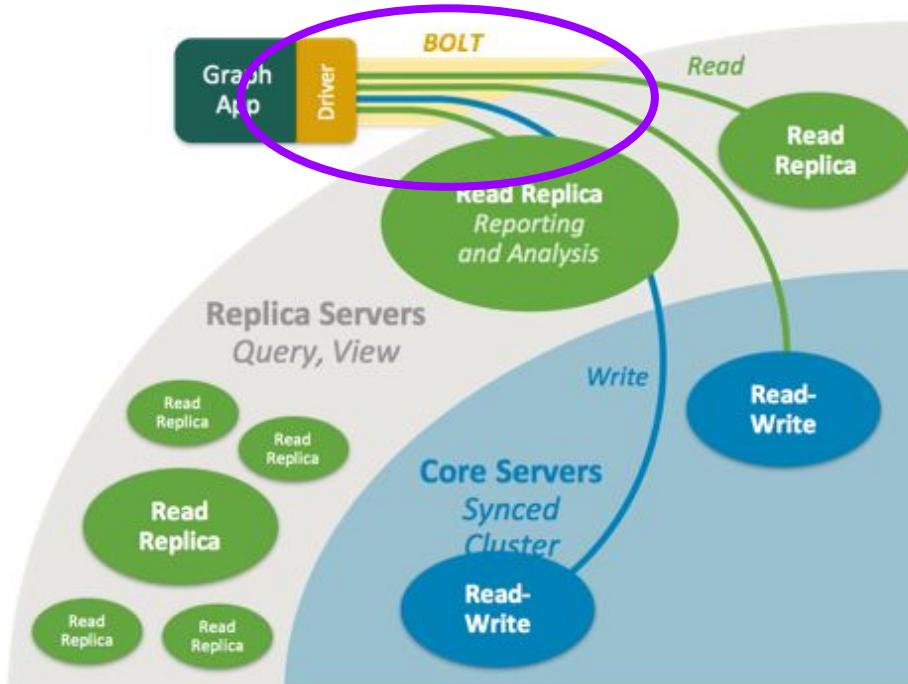
- **Community-Led Development:** Collaborative ecosystem for diverse use cases
- **Transparent Standardization:** Clear contribution processes and formal industry standardization

# **Developing: Driver Concepts:**

# Required Knowledge

- Basic Networking
- OOP
- Streams
- Basic High-Availability Concepts
- Common DB Paradigms

# Understanding this layer = [ Good Performance ]



# **Canonical \* Driver Model**



Uniform &  
Idiomatic

[nagels]

## All Neo4j Drivers

Thanks to the Neo4j contributor community, there are additionally drivers for almost every popular programming language, most of which mimic existing database driver idioms and approaches. Get started with your stack now, see the dedicated page for more detail.

<input checked="" type="checkbox"/> Java	<input checked="" type="checkbox"/> Spring	<input checked="" type="checkbox"/> Neo4j-OGM	<input checked="" type="checkbox"/> .NET	<input checked="" type="checkbox"/> JavaScript
<input checked="" type="checkbox"/> Python	<input checked="" type="checkbox"/> Go	Ruby	PHP	Erlang/Elixir
Perl	C/C++	Clojure	Haskell	R

If you are new to development, we recommend the [Jetbrains IDE ↗](#) for a good developer experience, which also comes with a [Neo4j Database plugin ↗](#).

<https://neo4j.com/developer/language-guides/#neo4j-drivers>

# Driver Concepts

- Driver
- Session
- Transaction
- Statement
- Result

# Driver



The **driver** is the thread-safe backbone providing access to Neo4j. It owns a **connection pool** through which all database connectivity occurs and can spawn **sessions** for carrying out work.

# Driver: Typical example



create\_driver.php

```
use Laudis\Neo4j\Basic\Driver;
use Laudis\Neo4j\Datasets\DriverConfiguration;

$driver = Driver::create(
    uri: 'neo4j://user:mypassword@localhost:7687',
    configuration: DriverConfiguration::create() →withUserAgent('MyApp/1.0.0')
);
```

# Driver: Connectivity

<scheme>://<user>:<password>@<host>:<port>

- Scheme determines the **behaviour** of the driver
- User and password is **optional**
- The host is the IP or Hostname to connect to **initially**
- The port is self explanatory

# Driver: Schema Connectivity

Certificate Type	Neo4j Causal Cluster	Neo4j Standalone Server	Direct Connection to Cluster Member
Unencrypted	neo4j	neo4j	bolt
Encrypted with Full Certificate	neo4j+s	neo4j+s	bolt+s
Encrypted with Self-Signed Certificate	neo4j+ssc	neo4j+ssc	bolt+ssc
Neo4j AuraDB ↗	neo4j+s	N/A	N/A

# Driver: Neo4j Scheme: Client side routing

1. Driver fetches the **routing table** from the original connection URI
1. The routing table is a **list of addresses** and their role
1. When creating a session the driver **lends a connection** based on the **session and server role**

# Driver: Configuration

Name	Type	Description
uri	String	Uri for initial connection
auth	Dictionary	Authentication information
user_agent	String	The user agent information
encrypted	Boolean	Enable encryption

# Driver: Authentication

Name	Data	Description
Basic	username + password	Default and basic, also works using URI
OIDC	token	Use an OpenID connect token
Kerberos	token	Use a Kerberos token
None	-	Access without a password, also works using URI

# Driver: Key Takeaways

- Excellent candidate **Singleton Pattern**
- **UX is the same** regardless of underlying protocol
- Neo4j schema is **default**
- Bolt schema trades **flexibility** for **performance**
- Connection URI is only for the **initial connection**
- A connection is against a Neo4j Instance, **not a database**
- The **connection pool** can be configured at this level

# Session



**Sessions** are lightweight containers for causally chained sequences of transactions. They borrow **connections** from the connection pool as required and chain transactions using **bookmarks**.

# Session: Typical Code Example



create\_session.php

```
use Laudis\Neo4j\Datasets\SessionConfiguration;
use Laudis\Neo4j\Enum\AccessMode;

$session = $driver->createSession(SessionConfiguration::create()
    ->withDatabase('my-database')
    ->withAccessMode(AccessMode::READ())
);
```

# Session: Configuration

Name	Type	Description
default_access_mode	String	READ or WRITE
database	Dictionary	The database the session works on
bookmarks	List<String>	Bookmarking
imp_user	String	Which user to impersonate
fetch_size	Integer	How many rows are fetched at once

# Session: Bookmarks: The Problem It Solves



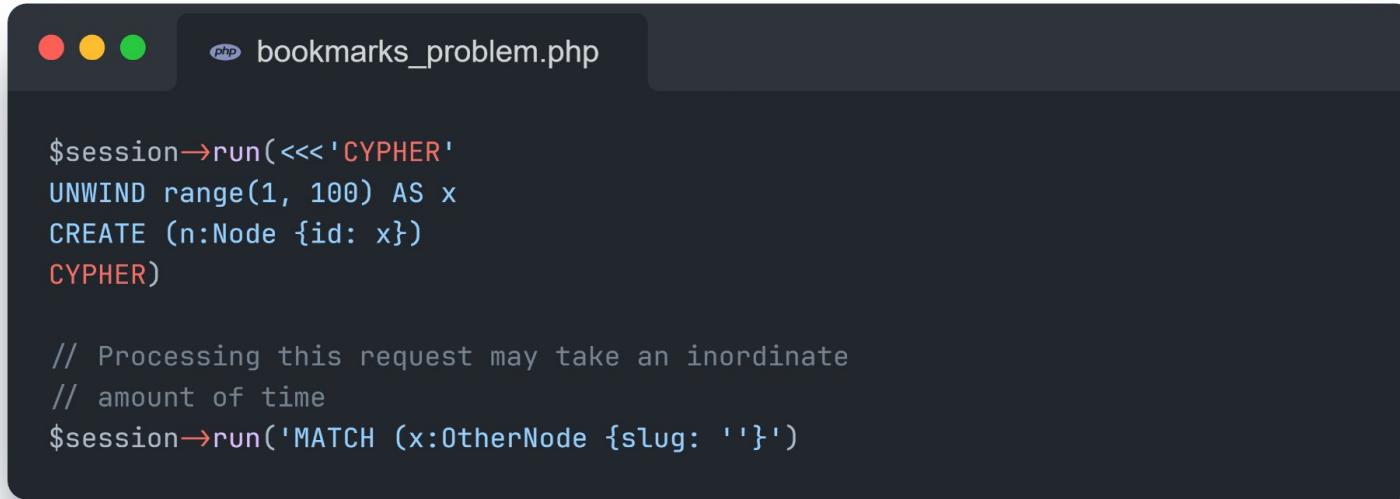
bookmarks\_problem.php

```
$session→run('CREATE (x:Test {name: "test"})');
// This may throw an exception
// as it may not find a node with name "test"
$session→run('MATCH (x:Test {name: "test"}) RETURN x')→first();
```

# Session: Bookmarks

- Neo4j provides a bookmark **identifier** marking the **future state** when the work is done
- The session **stores** this bookmark and **passes** it onto the next request
- Neo4j will only handle the new request once the **state of the database matches** the identifier

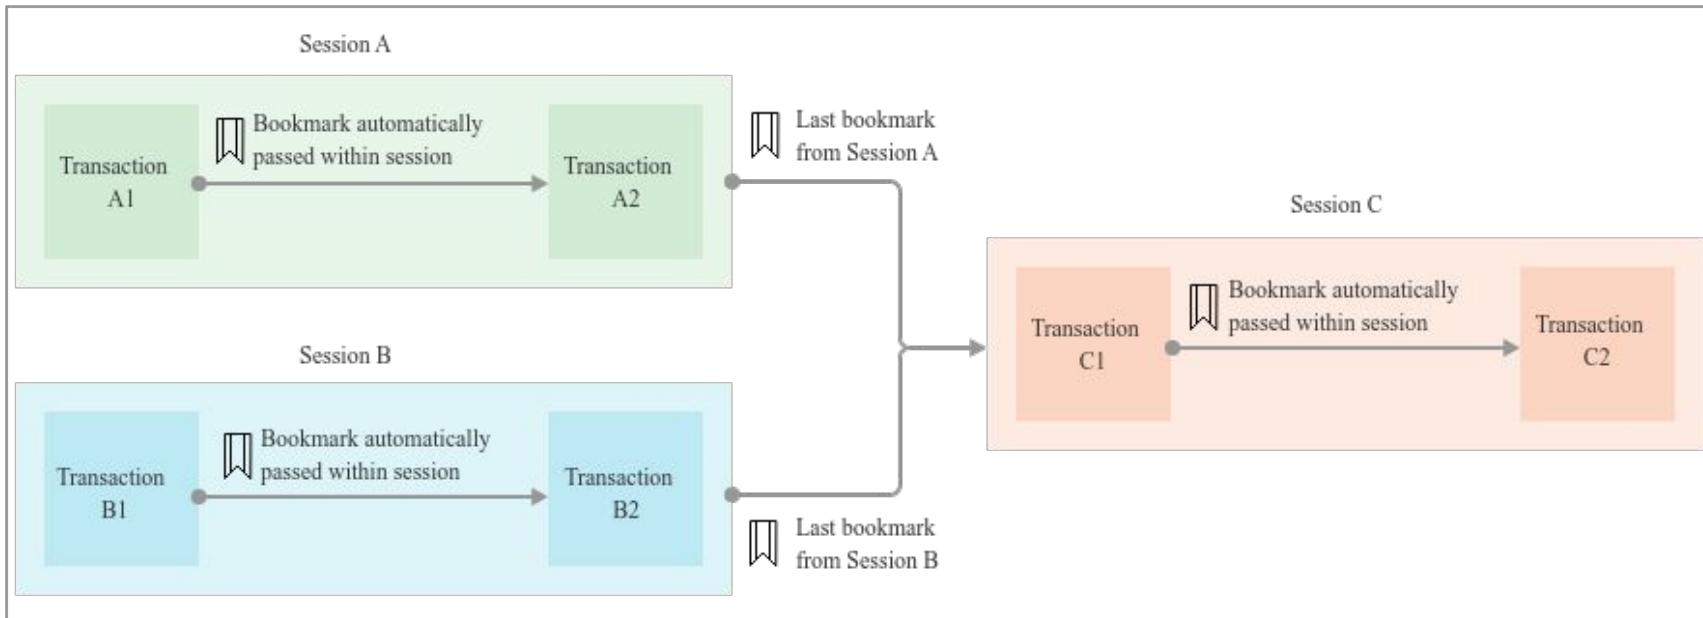
# Session: Bookmarks: The Problem It Creates



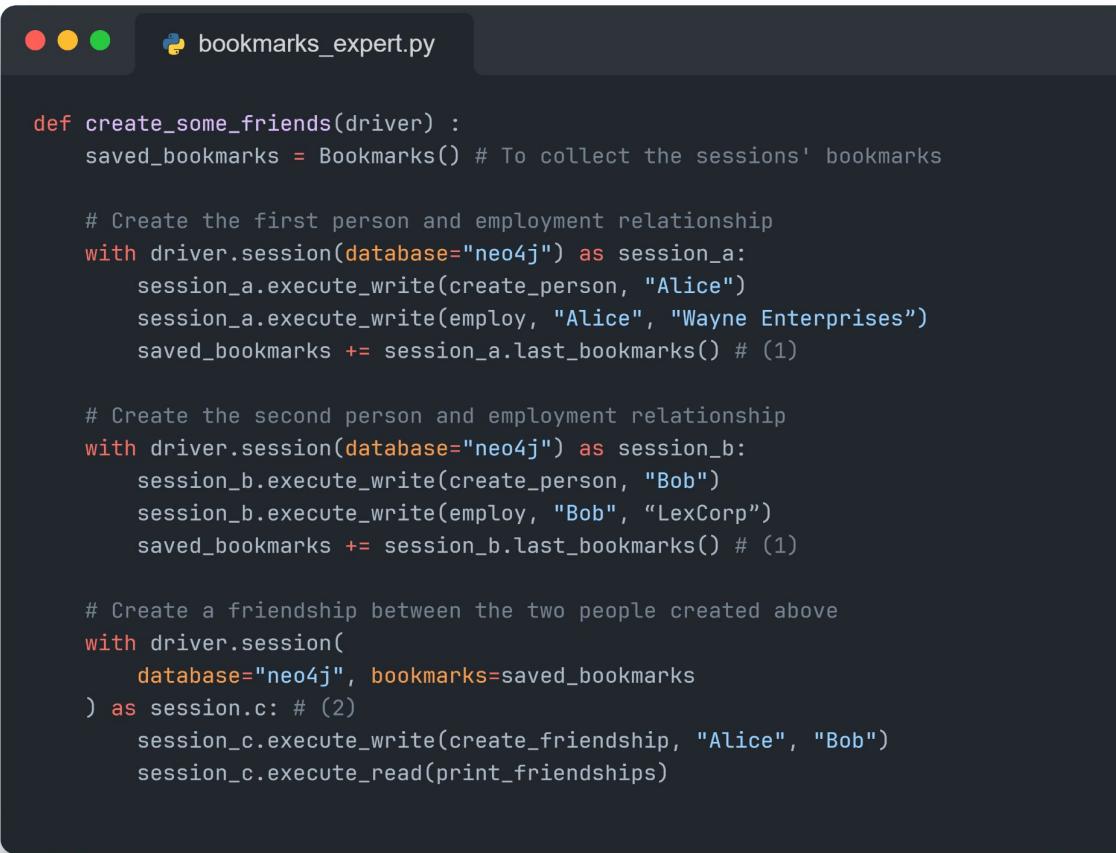
The screenshot shows a macOS application window titled "bookmarks\_problem.php". The window has the standard OS X title bar with red, yellow, and green buttons. The main content area displays the following PHP code:

```
$session→run(<<<'CYPHER'  
UNWIND range(1, 100) AS x  
CREATE (n:Node {id: x})  
CYPHER)  
  
// Processing this request may take an inordinate  
// amount of time  
$session→run('MATCH (x:OtherNode {slug: ''})')
```

# Session: Bookmarks: Expert Edition



# Session: Bookmarks: Expert Edition



```
bookmarks_expert.py

def create_some_friends(driver):
    saved_bookmarks = Bookmarks() # To collect the sessions' bookmarks

    # Create the first person and employment relationship
    with driver.session(database="neo4j") as session_a:
        session_a.execute_write(create_person, "Alice")
        session_a.execute_write(employ, "Alice", "Wayne Enterprises")
        saved_bookmarks += session_a.last_bookmarks() # (1)

    # Create the second person and employment relationship
    with driver.session(database="neo4j") as session_b:
        session_b.execute_write(create_person, "Bob")
        session_b.execute_write(employ, "Bob", "LexCorp")
        saved_bookmarks += session_b.last_bookmarks() # (1)

    # Create a friendship between the two people created above
    with driver.session(
        database="neo4j", bookmarks=saved_bookmarks
    ) as session_c: # (2)
        session_c.execute_write(create_friendship, "Alice", "Bob")
        session_c.execute_read(print_friendships)
```

# Session: Key Takeaways

- When in doubt, **create a new session**
- Many **open sessions** means many **open connections**
- Only reuse a session if the **work is related**
- **Don't worry** about bookmarks until you are optimising
- The session determines **the database** you work on
- Finetune result fetching and memory usage with **fetch\_size** config

# Transaction



**Transactions** are atomic units of work that may contain one or more query. Each transaction is bound to a single **connection** and is represented in the causal chain by a **bookmark**.

# Transaction: Configuration

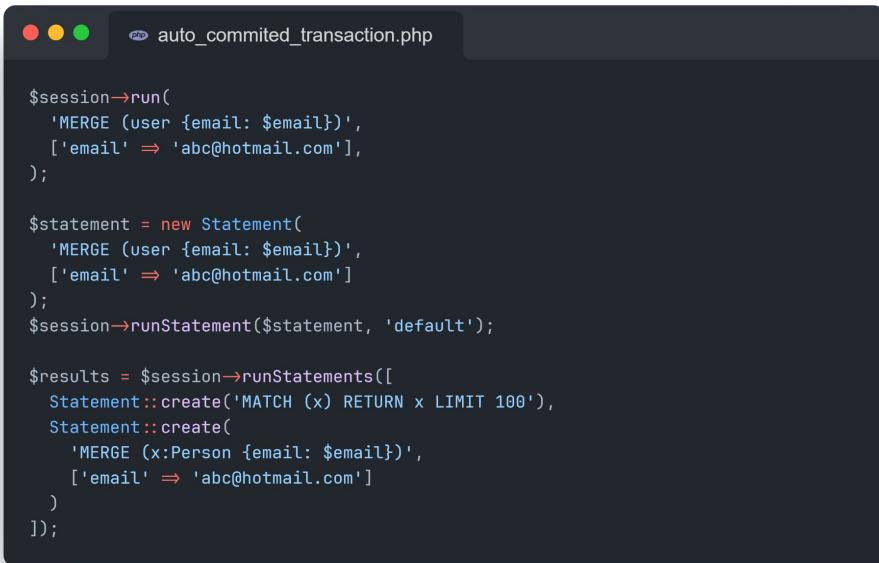
Name	Type	Description
metadata	Dictionary	Metadata to pass to Neo4j
timeout	Integer	The maximum duration of the transaction

# Transaction: Types

- **Auto-committed** queries
- **Transaction** functions
- **Unmanaged** transactions

# Transaction: Auto-committed

- Simple
- Intuitive
- May fail in complex scenarios



A screenshot of a Mac OS X-style application window titled "auto\_committed\_transaction.php". The window contains PHP code demonstrating an auto-committed transaction. The code uses the Doctrine DBAL library to run multiple database statements. It first runs a single MERGE statement, then creates a new Statement object for another MERGE, and finally runs both statements with a default transaction. The code also includes a query to match 100 rows.

```
$session->run(
    'MERGE (user {email: $email})',
    ['email' => 'abc@hotmail.com'],
);

$statement = new Statement(
    'MERGE (user {email: $email})',
    ['email' => 'abc@hotmail.com']
);
(session->runStatement($statement, 'default');

$results = $session->runStatements([
    Statement::create('MATCH (x) RETURN x LIMIT 100'),
    Statement::create(
        'MERGE (x:Person {email: $email})',
        ['email' => 'abc@hotmail.com']
    )
]);
```

# Transaction: Transient Errors

- **Database** related (memory, leader election change, ...)
- **Temporary**
- Requires a **retry**

# Transaction: Transaction Functions

- Manages Transient Errors for you
- Simple
- A bit more boilerplate
- Works in every scenario
- Need to be idempotent



```
use Laudis\Neo4j\Contracts\TransactionInterface;

// Do a simple merge and return the result
$result = $client->writeTransaction(static function (TransactionInterface $tsx) {
    $result = $tsx->run('MERGE (x {y: "z"}:X) return x');
    return $result->first()->get('x')['y'];
});

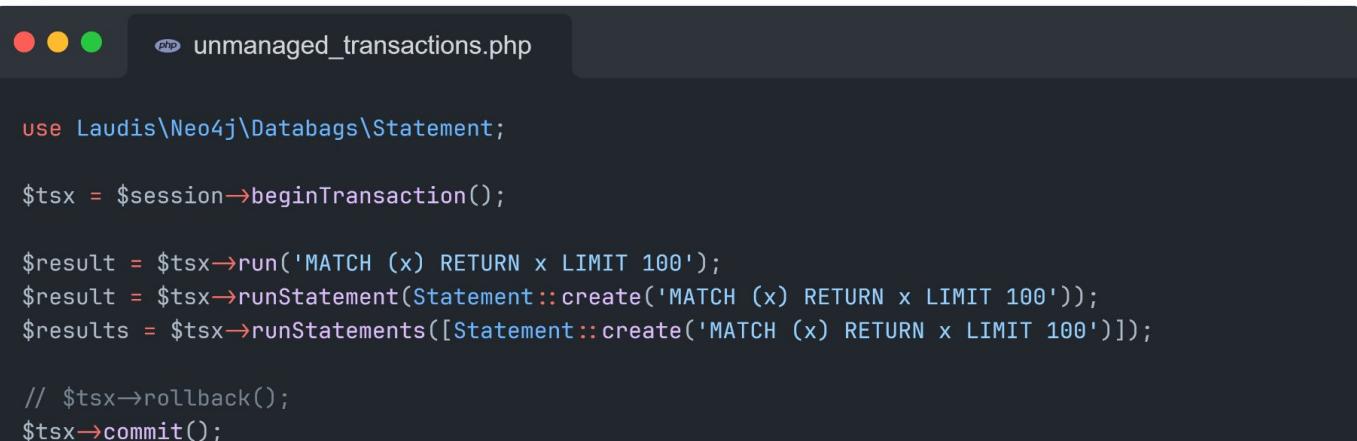
// Will result in an error
$client->readTransaction(static function (TransactionInterface $tsx) {
    $tsx->run('MERGE (x {y: "z"}:X) return x');
});

// This is a poorly designed transaction function
$client ->writeTransaction(static function (TransactionInterface $tsx) use ($externalCounter) {
    $externalCounter->incrementNodesCreated();
    $tsx->run('MERGE (x {y: $id}:X) return x', ['id' => Uuid::v4()]);
});

// This achieves the same effect but is safe in case it should be retried.
// The function is now idempotent.
$id = Uuid::v4();
$client->writeTransaction(static function (TransactionInterface $tsx) use ($id) {
    $tsx->run('MERGE (x {y: $id}:X) return x', ['id' => $id]);
});
$externalCounter->incrementNodesCreated();
```

# Transaction: Unmanaged

- Simple concept
- Complete control
- Punishes forgetfulness



A screenshot of a terminal window titled "unmanaged\_transactions.php". The window shows a block of PHP code. The code uses the `Laudis\Neo4j\Datasets\Statement` class to manage a transaction. It starts by creating a new transaction object, then runs three separate queries: one directly via `$tsx->run`, one via `$tsx->runStatement`, and one via `$tsx->runStatements`. Finally, it attempts to commit the transaction.

```
use Laudis\Neo4j\Datasets\Statement;

$tsx = $session->beginTransaction();

$result = $tsx->run('MATCH (x) RETURN x LIMIT 100');
$result = $tsx->runStatement(Statement::create('MATCH (x) RETURN x LIMIT 100'));
$results = $tsx->runStatements([Statement::create('MATCH (x) RETURN x LIMIT 100')]);

// $tsx->rollback();
$tsx->commit();
```

# Transaction: Key Takeaways

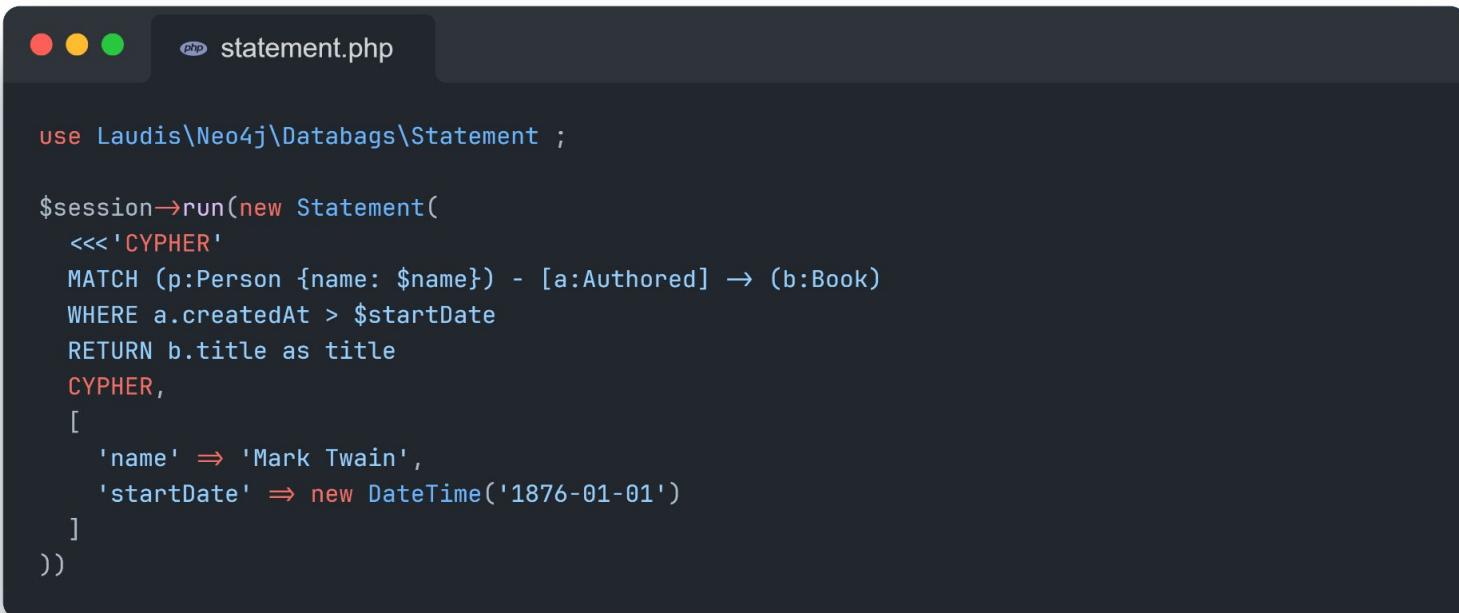
- **Functions** are the preferred way
- **Complicated** scenarios require **unmanaged** transactions
- **One connection** can have many **transactions** open
- Transactions **cannot be nested**
- Keep it **simple** by using **auto-committed** transactions

# Statement



A **statement** is a single instruction consisting of a **cypher** string and a **parameter** dictionary.

# Statement: Code Example



The screenshot shows a Mac OS X window titled "statement.php". The code inside the window is as follows:

```
use Laudis\Neo4j\Datasets\Statement ;

$session->run(new Statement(
    <<< 'CYPHER'
    MATCH (p:Person {name: $name}) - [a:Authored] -> (b:Book)
    WHERE a.createdAt > $startDate
    RETURN b.title as title
    CYPHER,
    [
        'name' => 'Mark Twain',
        'startDate' => new DateTime('1876-01-01')
    ]
))
```

# Statement Sidenote: Domain Specific Objects

Starred PHP rows can  
be used as a parameter

Cypher	PHP
null	* null
string	* string
integer	* int
float	* float
boolean	* bool
Map	* \Laudis\Neo4j\Types\CypherMap
List	* \Laudis\Neo4j\Types\CypherList
Point	* \Laudis\Neo4j\Contracts\PointInterface **
Date	* \Laudis\Neo4j\Types\Date
Time	* \Laudis\Neo4j\Types\Time
LocalTime	* \Laudis\Neo4j\Types\LocalTime
DateTime	* \Laudis\Neo4j\Types\DateTime
DateTimeZoneId	* \Laudis\Neo4j\Types\DateTimeZoneId
LocalDateTime	* \Laudis\Neo4j\Types\LocalDateTime
Duration	* \Laudis\Neo4j\Types\Duration
Node	\Laudis\Neo4j\Types\Node
Relationship	\Laudis\Neo4j\Types\Relationship
Path	\Laudis\Neo4j\Types\Path

# Statement: Key Takeaways

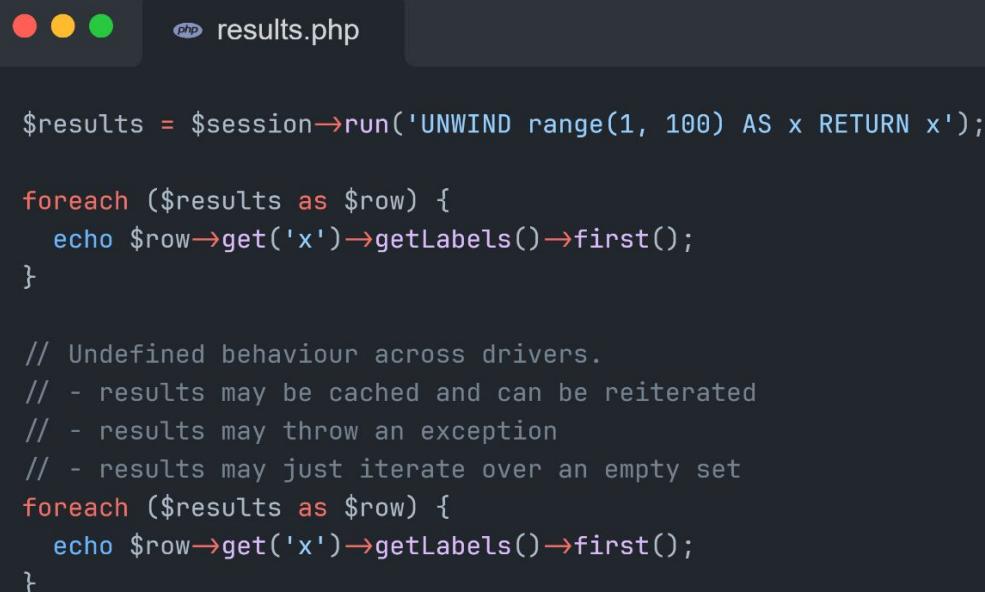
- Use **parameters everywhere**
- Parameters help with **query caching**
- Parameters stop **Cypher injections**
- Certain **objects** can be parameters
- Parameters are a **dictionary** whose **keys** directly map to their Cypher **parameter name**.

# Result



**Results** are the values returned by a query in the form of **rows**. Each result is bound to a single **connection** and can be **consumed** by iterating over it.

# Results: Cursors By Default



The screenshot shows a dark-themed macOS application window titled "results.php". Inside the window, there is a code editor displaying the following PHP code:

```
$results = $session->run('UNWIND range(1, 100) AS x RETURN x');

foreach ($results as $row) {
    echo $row->get('x')->getLabels()->first();
}

// Undefined behaviour across drivers.
// - results may be cached and can be reiterated
// - results may throw an exception
// - results may just iterate over an empty set
foreach ($results as $row) {
    echo $row->get('x')->getLabels()->first();
}
```

# Results: Metadata



result\_metadata.php

```
$results = $session->run(<<< 'CYPHER'  
UNWIND range(1, 100) AS x  
CREATE (n:Node {id: x})  
RETURN n  
CYPHER);  
  
echo $results->getSummary()->getCounters()->getNodesCreated(); // 100
```

# Results: Metadata Overview

## © ResultSummary

- ➊ counters: SummaryCounters
- ➋ databaseInfo: DatabaseInfo
- ➌ notifications: CypherList
- ➍ plan: Plan|null
- ➎ profiledPlan: ProfiledPlan|null
- ➏ statement: Statement
- ➐ queryType: QueryTypeEnum
- ➑ resultAvailableAfter: float
- ➒ resultConsumedAfter: float
- ➓ serverInfo: ServerInfo

# Results: Key Takeaways

- Results are **cursors**, read them once
- Results keep a lot of **useful metadata** for debugging or logging
- Results keep a **connection open** until they are consumed
- Results return neo4j **specific objects** too
- Querying some of the results metadata requires **full consumption** of the results

# Final Takeaways

- Understand the building blocks
- Optimise by using lower level configurations carefully
- Debug and Profile with result summaries
- Graph Databases support real time applications
- Graph Databases have great ecosystems for data science
- Have fun!

[nagels]

Rolling Your own MCP Server

# DEMO TIME

Copyright © 2025 Nagels



The end

# Questions?

Want this slide deck for free?  
Contact Me

