

Unlocking the Power of Graph Databases with Neo4j

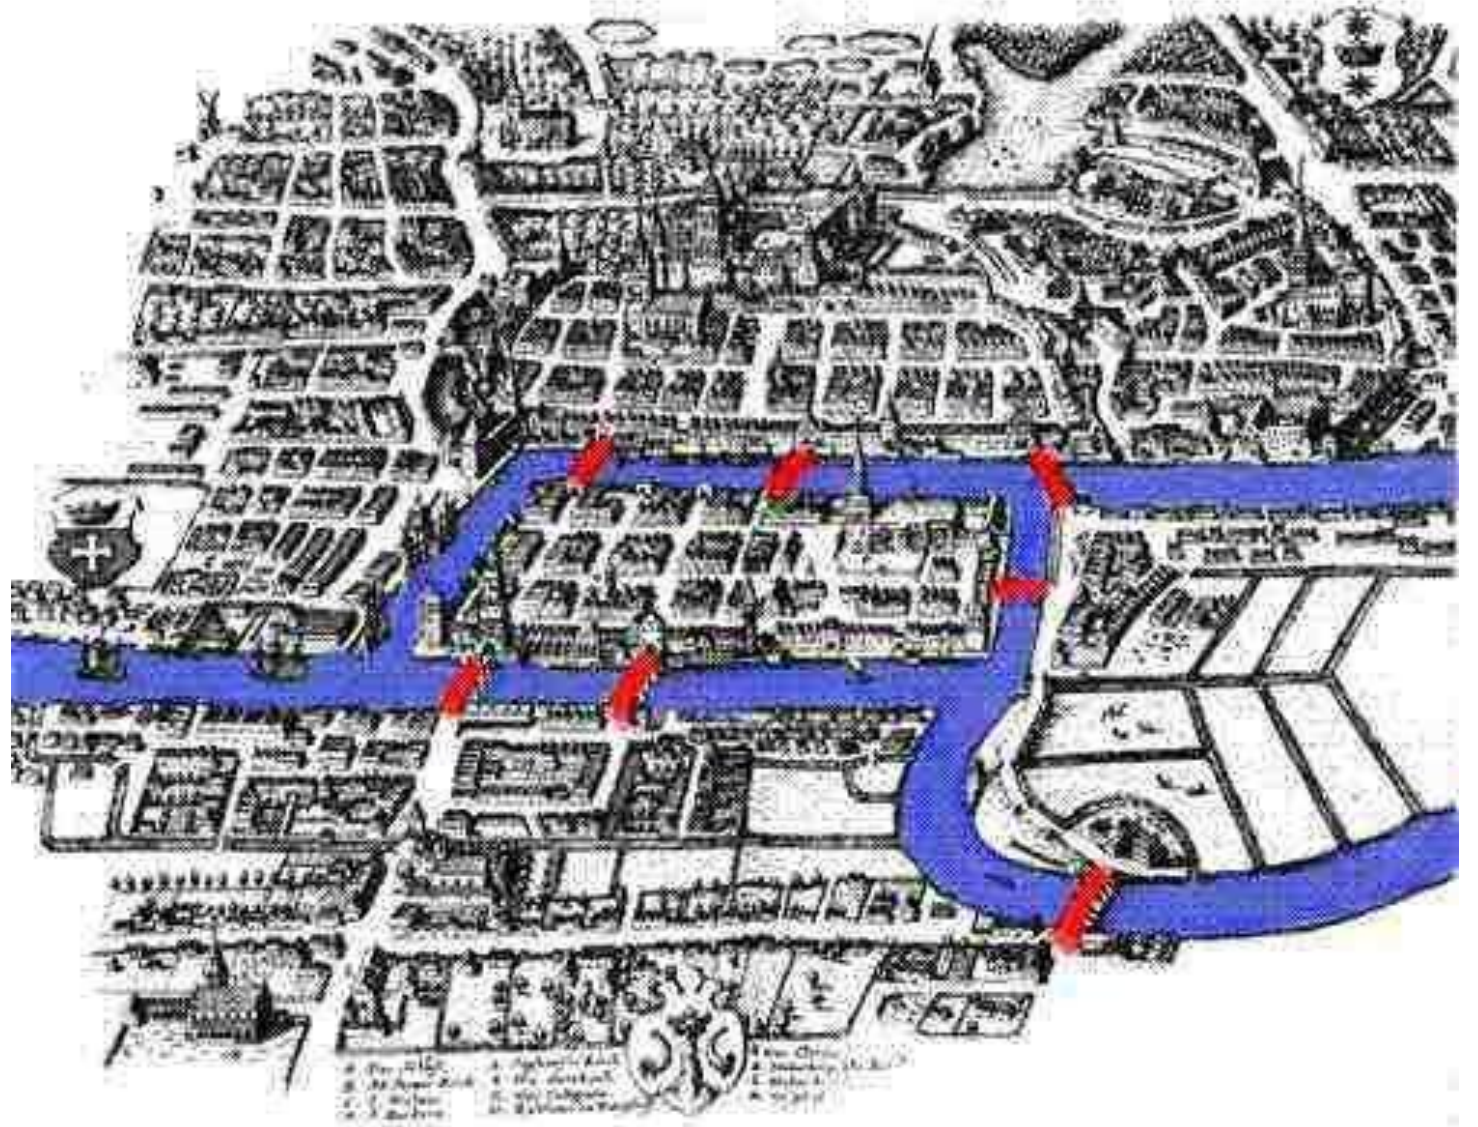
Ilya Verbitskiy

Agenda

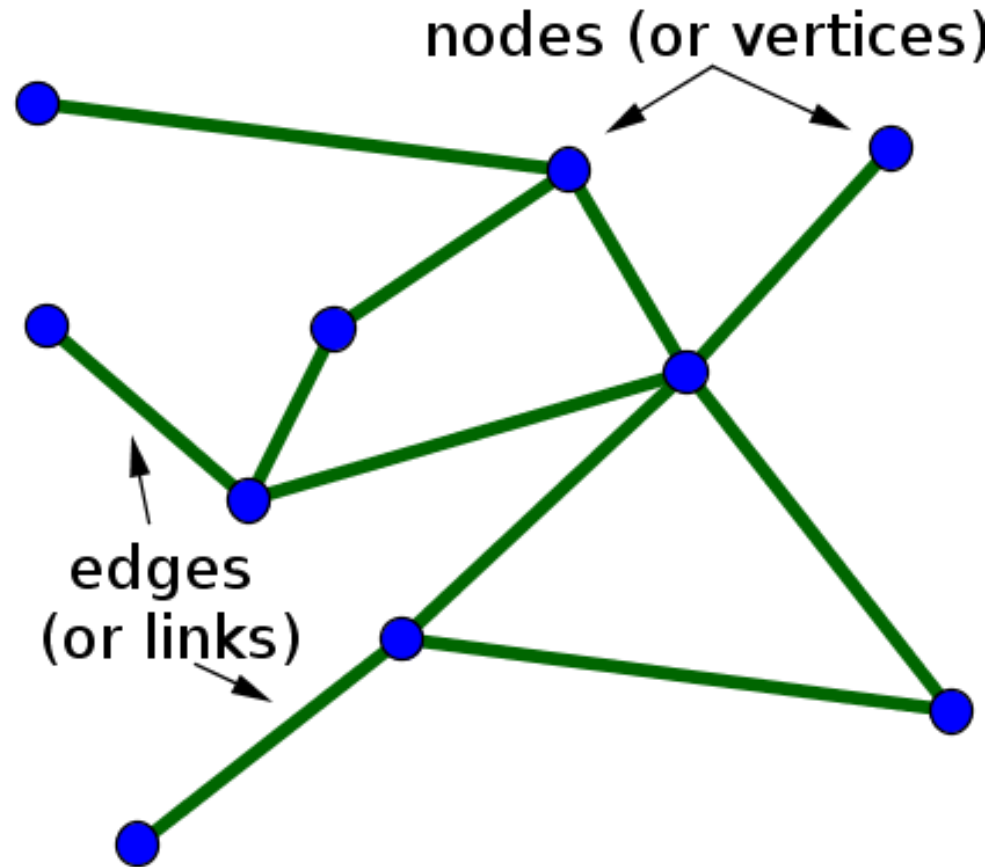
- Introduction to graph theory and Neo4j
- Data Modelling - Northwind
- Setting up the Northwind database on Neo4j AuraDB
- Introduction to Cypher

Königsberg, Prussia, 1736

- Is it possible for us to stroll through the city in such a way that we traverse each of the seven bridges just one time?

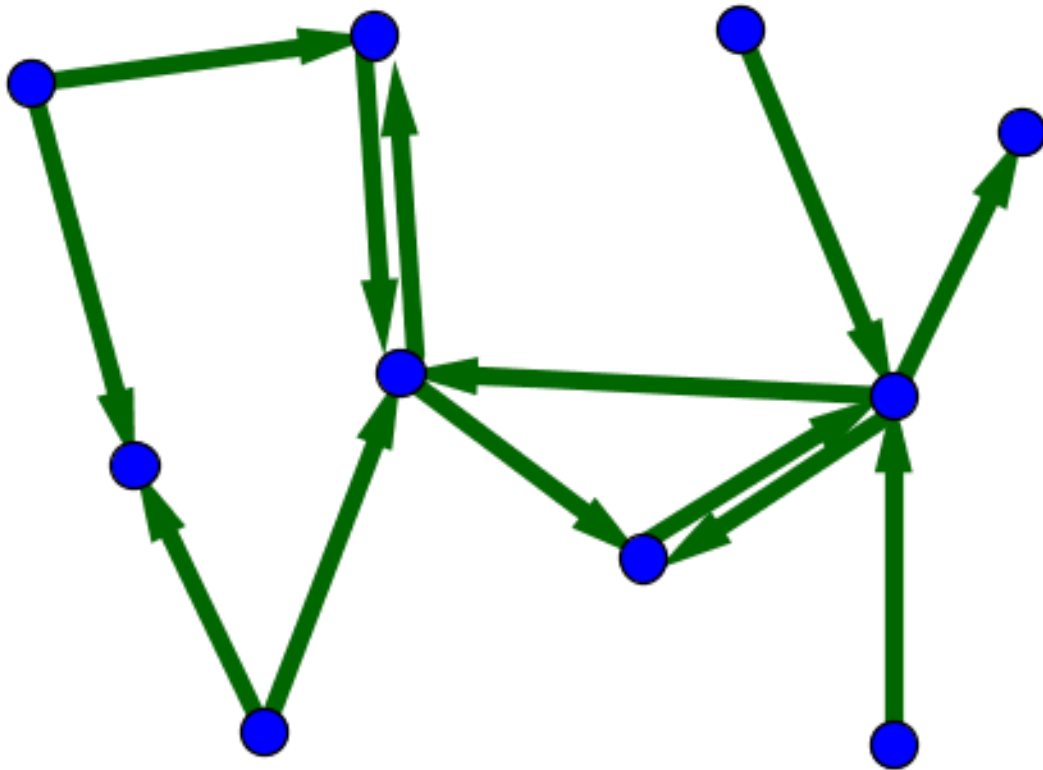


What is a graph

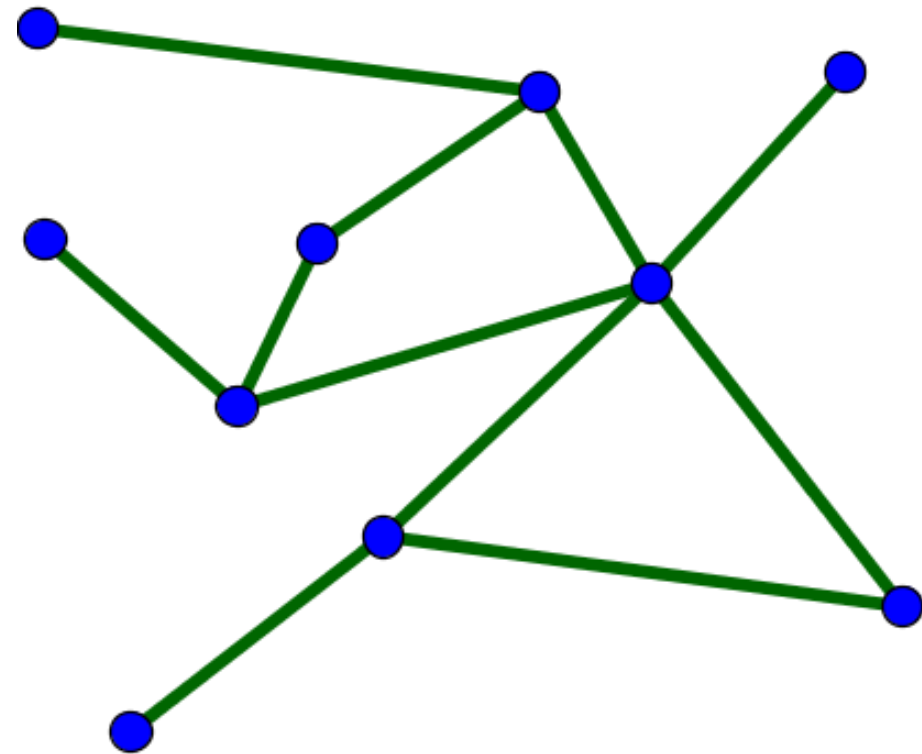


Node relationships

A directed network



An undirected network



Graphs around us

- Logistics
- Finance
- Network and IT operations
- E-Commerce
- Crime investigations
- Generative AI

Why Graph database

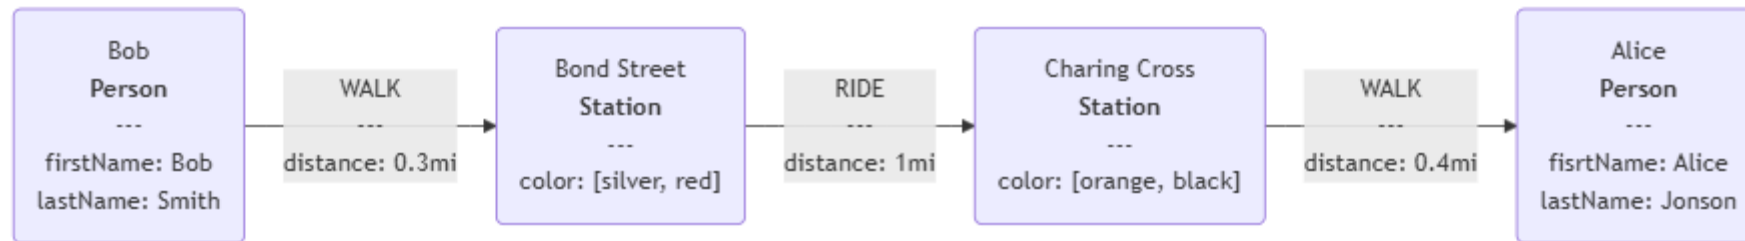
- Navigate deep hierarchies
- Find hidden connections between distant items
- Discover inter-relationships between items

What is Neo4j

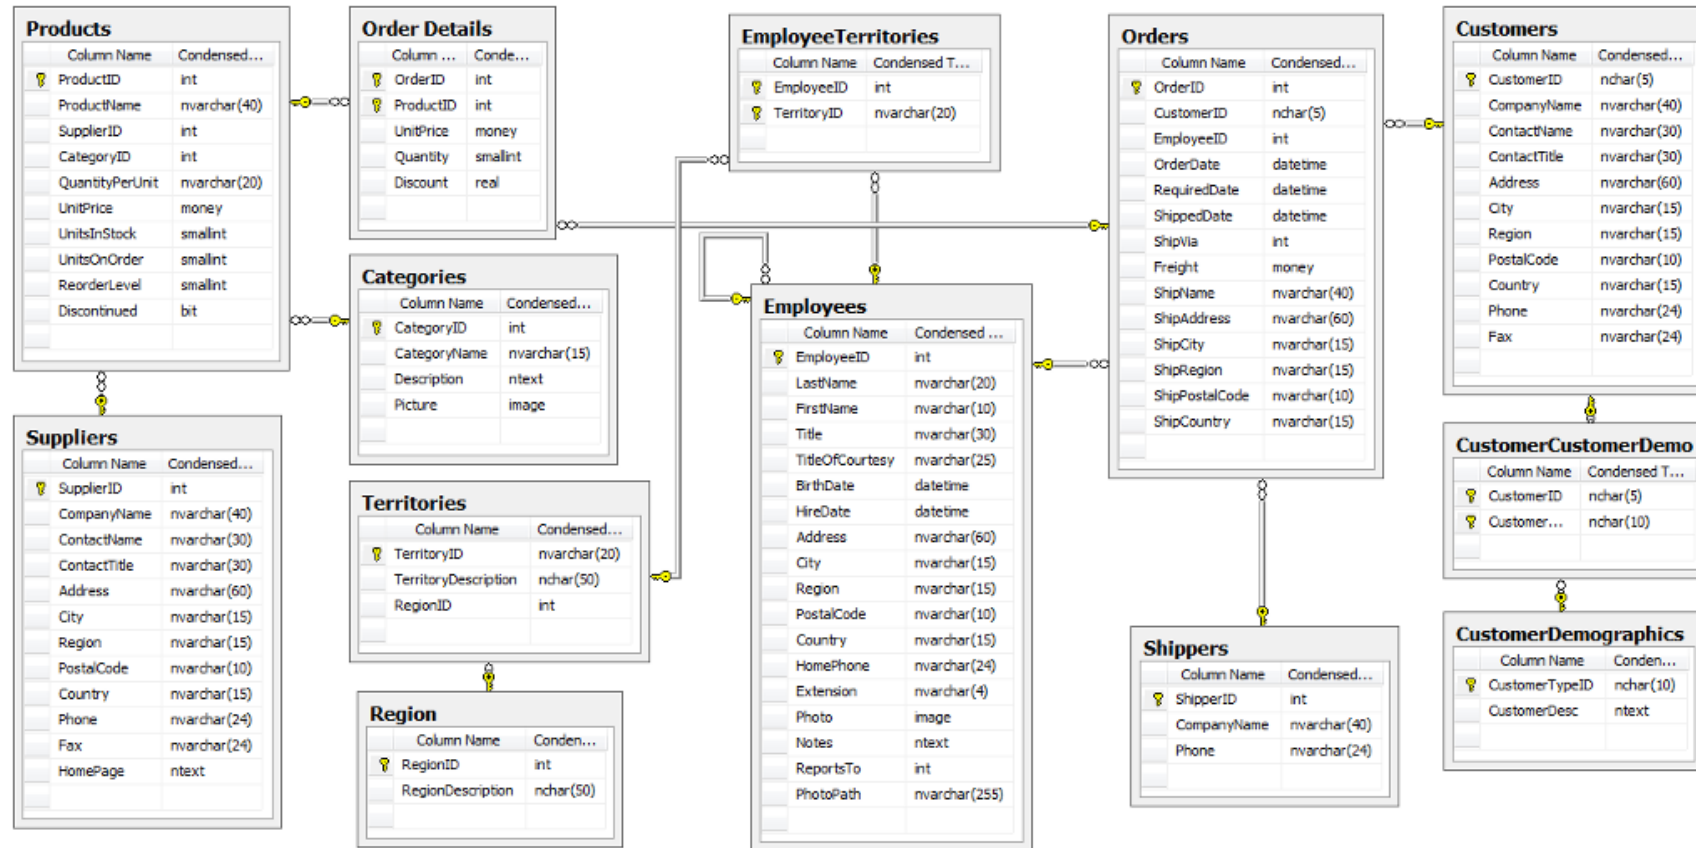
- Native graph database designed specifically with traversal in mind
- High performance read and write scalability
- ACID compliant
- Index-free adjacency (IFA)
- Cypher query language
- Open Source (GPLv3)
- Cloud enabled (AuraDB)

Neo4j data elements

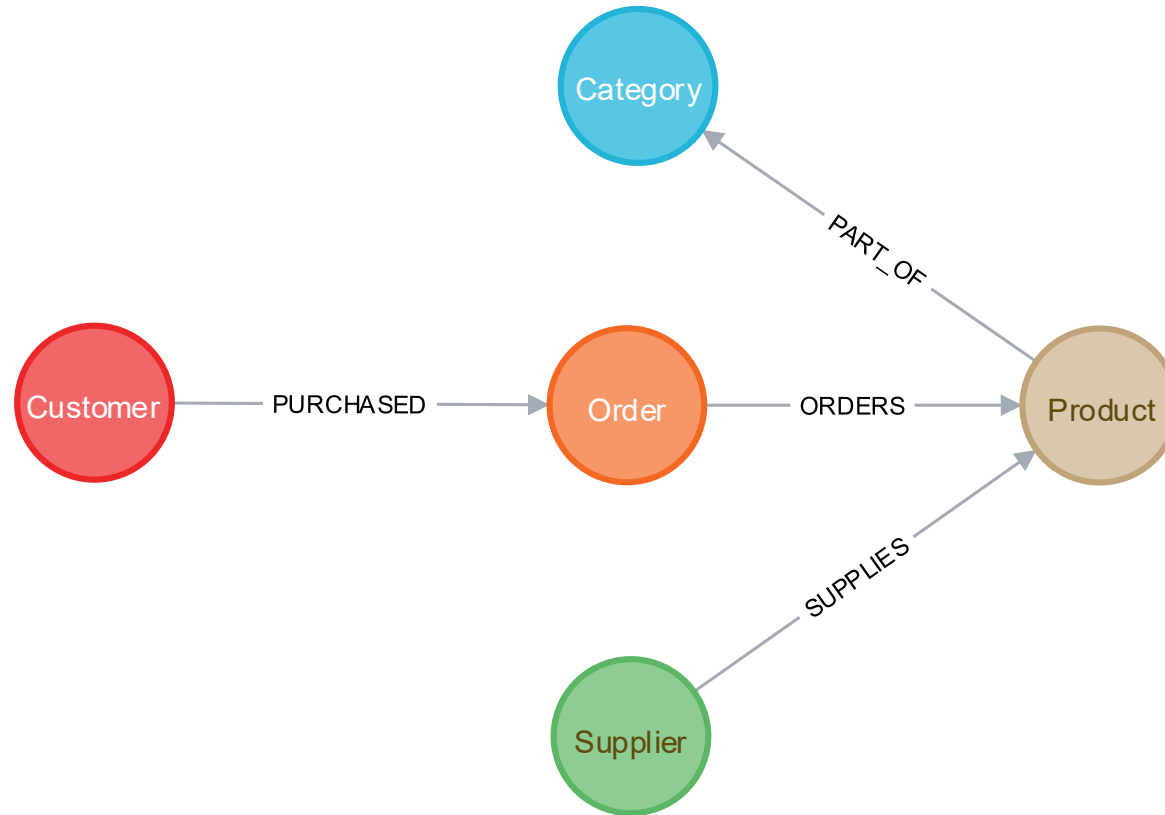
- Nodes
- Relationships
- Labels
- Properties



Data Modeling – Northwind (SQL)



Data Modeling – Northwind (Neo4j)



<https://github.com/neo4j-graph-examples/northwind>



Setting up Northwind database

neo4j

aura

WebStoating / Demo

Send feedback

Learn

Ilya Verbitskiy

Data services

Instances

Import

Data APIs

Beta

Tools

Query

Explore

Operations

Metrics

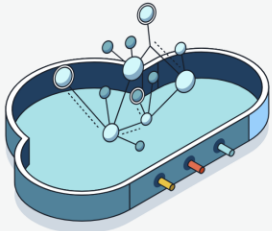
Logs

Project

Users

Billing

Settings



Create your first instance

Configure your instance and start realizing the endless possibilities with graphs.

Create instance

Have data? [Import](#) existing data into Neo4j.

Setting up Northwind database

The screenshot displays the Neo4j Aura web interface. The top navigation bar includes the Neo4j logo, the text "WebStoating / Demo", and links for "Send feedback", "Learn", and a user profile for "Ilya Verbitskiy". Below the navigation bar, the interface is divided into three main sections:

- Data services:** A sidebar on the left containing links for "Instances", "Import", "Data APIs", "Tools", "Operations", and "Project". The "Tools" section is expanded, showing "Query" (highlighted), "Explore", "Metrics", "Logs", "Users", "Billing", and "Settings".
- Database information:** A central panel showing details for the selected database. It includes a "Database information" header, a "Nodes (0)" section with a plus icon, a "Relationships (0)" section with a plus icon, and a "Property keys" section.
- Query editor:** A large panel on the right for writing and executing queries. It features a query input field with the text "neo4j\$" and a "Query" button. Below the input field, there are two main sections: "GUIDE" and "DATASET".

The "GUIDE" section is titled "Query fundamentals" and lists the following topics you'll learn to:

- write basic queries;
- view graph and tabular results;
- perform queries to answer questions;
- perform advanced queries.

The "DATASET" section is titled "Try Neo4j with the Movie Graph" and describes a fully-built graph example to kickstart your Neo4j journey. It mentions discovering key query patterns through a fun, familiar world of actors and films.

At the bottom of the interface, there is a "Go to Learn Center" button with a graph icon.

Last update: 17:36:52

Setting up Northwind database

The screenshot displays the Neo4j Aura web interface. The top navigation bar includes the Neo4j logo, the text "WebStoating / Demo", and links for "Send feedback", "Learn", and a user profile for "Ilya Verbitskiy". The left sidebar is organized into sections: "Data services" (with "Instances" selected), "Tools" (including "Import", "Data APIs", "Query", and "Explore"), "Operations" (including "Metrics" and "Logs"), and "Project" (including "Users", "Billing", and "Settings"). The main content area shows the details for "Instance01", which is in a "RUNNING" state. It displays the instance ID "ff53a06f" and its configuration: "Type: AuraDB Free", "Nodes: 0 (0%)", and "Relationships: 0 (0%)". A "Metrics" link is visible below the configuration. A context menu is open over the instance details, offering actions such as "Inspect", "Snapshots", "Backup & restore", "Clone to", "Reset to blank", and "Delete". A "Create instance" button is located in the top right corner of the main content area.

neo4j aura WebStoating / Demo

Send feedback Learn Ilya Verbitskiy

Data services

- Instances
- Import
- Data APIs Beta

Tools

- Query
- Explore

Operations

- Metrics
- Logs

Project

- Users
- Billing
- Settings

Search

Create instance

Instance01

● RUNNING

ID: ff53a06f

Type: AuraDB Free Nodes: 0 (0%) Relationships: 0 (0%)

> Metrics

- Inspect
- Snapshots
- Backup & restore
- Clone to
- Reset to blank
- Delete

Setting up Northwind database

The screenshot displays the Neo4j Aura web interface. On the left, a sidebar contains navigation links for 'Data services' (Instances, Import, Data APIs), 'Tools' (Query, Explore), 'Operations' (Metrics, Logs), and 'Project' (Users, Billing, Settings). The main area shows 'Instance01' in a 'RUNNING' state with ID 'ff53a06f'. Below this, a 'Metrics' section is visible. On the right, a modal window titled 'Instance01' is open, showing the 'Restore from backup file' tab. This tab features a large dashed box with a cloud upload icon and the text: 'Drag & Drop your Neo4j .dump or .backup or .tar file here (4GB or less) or [browse](#)'. Below this, a note states: 'Want to upload a dump or backup larger than 4GB? Click [here](#) to view the Neo4j Admin 'upload' command in the Aura Docs. The graph you upload should be smaller than 200,000 nodes and 400,000 relationships.'

neo4j aura WebStoating / Demo

Send feedback ? Learn Ilya Verbitskiy

Data services

- Instances
- Import
- Data APIs Beta

Tools

- Query
- Explore

Operations

- Metrics
- Logs

Project

- Users
- Billing
- Settings

Search

Instance01 ● RUNNING

ID: ff53a06f

Type: AuraDB Free Nodes: 0 (0%) Relationships: 0 (0%)

> Metrics

Instance01 ● RUNNING

Overview Snapshots Restore from backup file

Drag & Drop your Neo4j .dump or .backup or .tar file here (4GB or less)

or [browse](#)

Want to upload a dump or backup larger than 4GB? Click [here](#) to view the Neo4j Admin 'upload' command in the Aura Docs.

The graph you upload should be smaller than 200,000 nodes and 400,000 relationships.

Setting up Northwind database

The screenshot displays the Neo4j Aura web interface. The top navigation bar includes the Neo4j logo, the text "WebStoating / Demo", and links for "Send feedback", "Learn", and a user profile for "Ilya Verbitskiy". The left sidebar is organized into sections: "Data services" (with "Instances" selected), "Tools" (including "Import", "Data APIs", "Query", and "Explore"), "Operations" (including "Metrics" and "Logs"), and "Project" (including "Users", "Billing", and "Settings"). The main content area features a search bar and a list of instances. "Instance01" is highlighted, showing its ID as "ff53a06f" and its status as "LOADING". Below this, it specifies the instance type as "AuraDB Free" and shows that there are 0 nodes and 0 relationships. A "Metrics" link is also visible. On the right, a detailed view for "Instance01" is open, showing tabs for "Overview", "Snapshots", and "Restore from backup file". A light blue notification box states: "Loading in progress. Import feature will be available again once current data loading completes."

Setting up Northwind database

The screenshot displays the Neo4j Aura web interface. The top navigation bar includes the Neo4j logo, the text "WebStoating / Demo", and links for "Send feedback", "Learn", and a user profile for "Ilya Verbitskiy". The left sidebar contains a menu with sections: "Data services" (Instances, Import, Data APIs), "Tools" (Query, Explore), and "Operations" (Metrics, Logs). The "Query" tool is selected, showing a "Database information" panel on the left and a query editor on the right.

Database information

Instance: Instance01 Database: neo4j User: Aura (iverbitskiy@gmail.com)

Nodes (1,035)

- Category
- Customer
- Order
- Product
- Supplier

Relationships (3,139)

- ORDERS
- PART_OF
- PURCHASED
- SUPPLIES

Property keys

- address
- categoryID
- categoryName
- city
- companyName
- contactName
- contactTitle
- country
- customerID
- description
- discontinued
- discount
- employeeID
- fax
- freight
- homePage
- orderDate
- orderID
- phone
- picture

[Show all \(20 more\)](#)

Last update: 18:12:53

Query editor

neo4j\$

\$:welcome

GUIDE

Query fundamentals

You'll learn to:

- write basic queries;
- view graph and tabular results;
- perform queries to answer questions;
- perform advanced queries.

DATASET

Try Neo4j with the Movie Graph

Dive into a fully-built graph example to kickstart your Neo4j journey. Discover key query patterns through a fun, familiar world of actors and films.

Go to Learn Center

Cypher basics

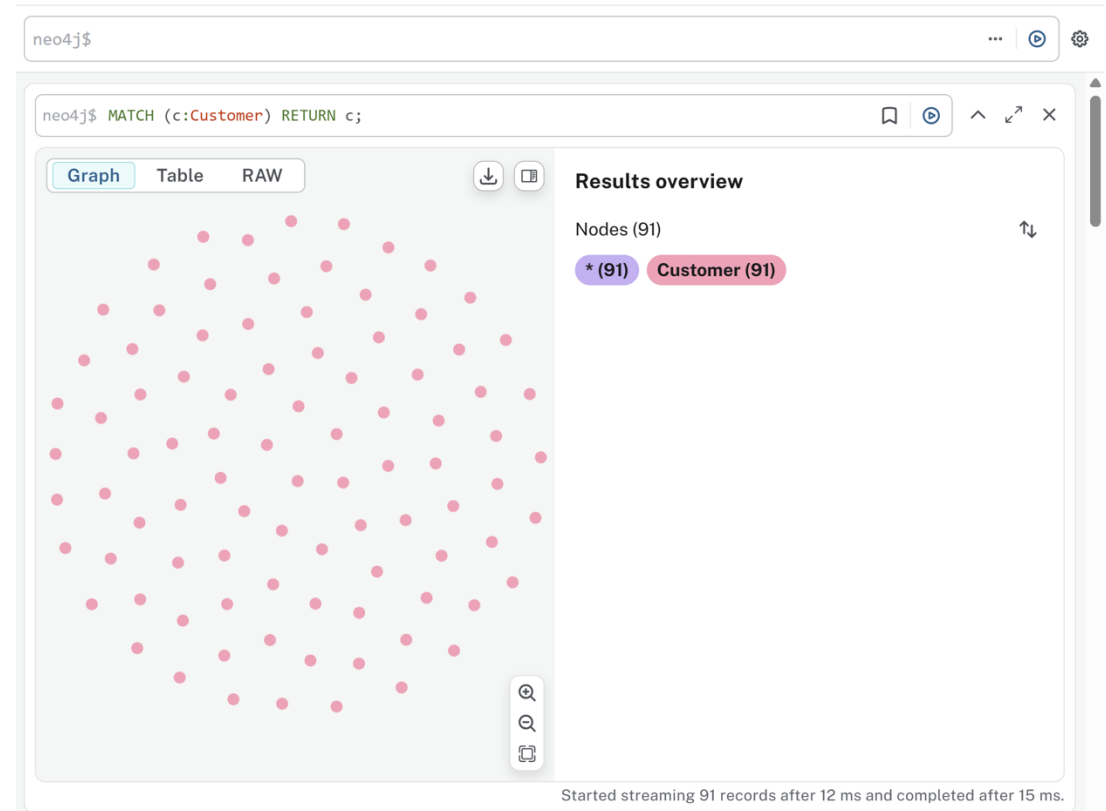
- **Nodes** are represented by parentheses ().
- **Labels** are specified using a colon :, e.g., (:Category).
- **Relationships** between nodes use two dashes --, e.g., (:Order)--(:Product).
- **Direction of relationships** is indicated with < or >, e.g., (:Order)-->(:Product).
- **Relationship types** are enclosed in square brackets [], e.g., [:ORDERS].
- **Properties** are key-value pairs, e.g., {employeeID: 4}.

SELECT #1

Get all columns from the Customers table.

SELECT * FROM Customers

MATCH (c:Customer) RETURN c



SELECT #2

Get the top 25 Customers alphabetically by Country and name.

```
SELECT TOP 25 *  
FROM Customers  
ORDER BY Country, ContactName  
  
MATCH (c:Customer)  
ORDER BY c.country, c.contactName  
RETURN c  
LIMIT 25;
```

The screenshot shows a database query interface with a query editor at the top and a results table below. The query is a Cypher query designed to retrieve the top 25 customers, ordered by country and then by contact name. The results are displayed in a table format with columns for the customer ID, the customer object, and the contact name. The first few rows show customers from Argentina, Austria, Belgium, and Brazil.

```
1 MATCH (c:Customer)  
2 ORDER BY c.country, c.contactName  
3 RETURN c  
4 LIMIT 25;
```

	c	
1	(:Customer {country: "Argentina", contactTitle: "Sales Agent", address: "Cerrito 333", p...	
2	(:Customer {country: "Argentina", contactTitle: "Sales Representative", address: "Av. de	
3	(:Customer {country: "Argentina", contactTitle: "Sales Agent", address: "Ing. Gustavo Mo	
4	(:Customer {country: "Austria", contactTitle: "Sales Manager", address: "Geislweg 14", p	
5	(:Customer {country: "Austria", contactTitle: "Sales Manager", address: "Kirchgasse 6",	
6	(:Customer {country: "Belgium", contactTitle: "Sales Agent", address: "Rue Joseph-Bens 5	
7	(:Customer {country: "Belgium", contactTitle: "Accounting Manager", address: "Boulevard	
8	(:Customer {country: "Brazil", contactTitle: "Sales Representative", address: "Av. Inês	
9	(:Customer {country: "Brazil", contactTitle: "Sales Associate", address: "Av. Brasil, 44	

Started streaming 25 records after 72 ms and completed after 80 ms.

SELECT #3

Get the count of all Orders made during 1997.

```
SELECT COUNT(*)  
FROM Orders  
WHERE YEAR(OrderDate) = 1997
```

```
MATCH (o:Order)  
WHERE o.orderDate.year = 1997  
RETURN COUNT(o);
```

SELECT #4

Get all orders placed on the 19th of May, 1997.

```
SELECT *
```

```
FROM Orders
```

```
WHERE OrderDate = '19970319'
```

```
MATCH (o:Order)
```

```
WHERE o.orderDate = date('1997-03-19')
```

```
RETURN o;
```

SELECT #5

Create a report for all the orders of 1996 and their Customers.

```
SELECT *  
FROM Orders o  
INNER JOIN Customers c ON o.CustomerID = c.CustomerID  
WHERE YEAR(o.OrderDate) = 1996
```

```
MATCH p=(c:Customer)-[:PURCHASED]->(o:Order)  
WHERE o.orderDate.year = 1996  
RETURN p;
```


SELECT #6

Create a report for all 1996 orders and their Customers. Return only the Order ID, Order Date, Customer ID, Name, and Country.

```
SELECT o.OrderID, o.OrderDate, c.CustomerID, c.ContactName, c.Country
FROM Orders o
INNER JOIN Customers c ON o.CustomerID = c.CustomerID
WHERE YEAR(o.OrderDate) = 1996
```

```
MATCH (c:Customer)-[:PURCHASED]->(o:Order)
WHERE o.orderDate.year = 1996
RETURN o.orderID, o.orderDate, c.customerID, c.contactName, c.country;
```

SELECT #7

Create a report that shows the number of customers from each city.

```
SELECT c.City, COUNT(*)  
FROM Orders o  
INNER JOIN Customers c ON o.CustomerID = c.CustomerID  
GROUP BY c.City
```

```
MATCH (c:Customer)-[:PURCHASED]->(o:Order)  
RETURN c.city, COUNT(*);
```

SELECT #8

Create a report that shows the total quantity of products ordered. Only show records for products for which the quantity ordered is fewer than 200

```
SELECT p.ProductName, SUM(od.Quantity) as Quantity
FROM OrderDetails od
INNER JOIN Products p ON od.ProductID = p.ProductID
GROUP BY p.ProductName
HAVING SUM(od.Quantity) < 200
ORDER BY Quantity
```

SELECT #8

```
MATCH (o:Order)-[od:ORDERS]->(p:Product)
WITH p.productName as productName, SUM(od.quantity) as quantity
WHERE quantity < 200
ORDER BY quantity
RETURN productName, quantity;
```

WITH Clause

- WITH Clause allows chaining query parts, passing results to the next step.
- Scope Control: Only variables included in WITH are carried forward; others are discarded.
- Wildcard * can be used to include all currently scoped variables.
- Data Manipulation: WITH enables modifying output before passing it to the following query part.
- Flexible Results: You can reshape data or filter the number of entries in the result set.

Aggregating functions

1. Counting and Grouping

- COUNT(*) – Counts all records.
- COUNT(expression) – Counts non-null values of an expression.
- COLLECT(expression) – Gathers values into a list.

2. Statistical Aggregates

- AVG(expression) – Calculates the average of numeric values.
- SUM(expression) – Computes the sum of numeric values.
- MIN(expression) – Returns the minimum value.
- MAX(expression) – Returns the maximum value.
- STDEV(expression) – Computes the standard deviation.
- STDEVP(expression) – Computes the population standard deviation.
- PERCENTILE_CONT(expression, percentile) – Returns the continuous percentile.
- PERCENTILE_DISC(expression, percentile) – Returns the discrete percentile.

SELECT #9

Create a report that shows the total number of orders by Customer since December 31, 1996. The report should only return rows for which the total number of orders is greater than 15

```
SELECT c.ContactName, COUNT(o.OrderID) as TotalOrders
FROM Orders o
INNER JOIN Customers c ON o.CustomerID = c.CustomerID
WHERE OrderDate > '1996-12-31'
GROUP BY c.ContactName
HAVING COUNT(o.OrderID) > 15
```

SELECT #9

```
MATCH (c:Customer)-[:PURCHASED]->(o:Order)
WHERE o.orderDate > date("1996-12-31")
WITH c.contactName as contactName, COUNT(o.orderID) as totalOrders
WHERE totalOrders > 15
RETURN contactName, totalOrders;
```


SELECT #10

Which UK Customers have paid more than 1000 dollars

```
SELECT c.ContactName, SUM(od.UnitPrice * od.Quantity * (1 - od.Discount)) as Paid
FROM Customers c
INNER JOIN Orders o ON c.CustomerID = o.CustomerID
INNER JOIN OrderDetails od ON o.OrderID = od.OrderID
WHERE c.Country = 'UK'
GROUP BY c.ContactName
HAVING SUM(od.UnitPrice * od.Quantity * (1 - od.Discount)) > 1000
```

SELECT #10

MATCH (c:Customer)-[:PURCHASED]->(o:Order)-[od:ORDERS]->(p:Product)

WHERE c.country = 'UK'

WITH c.contactName as contactName, SUM(toFloat(od.unitPrice) * od.quantity * (1 - toFloat(od.discount))) as paid

WHERE paid > 1000

RETURN contactName, paid;

INSERT/UPSERT #1

Insert yourself into the Customers table Include the following fields: CustomerID, CompanyName, ContactName, ContactTitle, Address, City, Region, PostalCode, Country, Phone, Fax

```
INSERT INTO Customers (CustomerID, CompanyName, ContactName, ContactTitle,  
Address, City, Region, PostalCode, Country, Phone, Fax)
```

```
VALUES
```

```
('ILYA1', 'Acme Corp', 'Ilya Verbitskiy', 'Manager', '123 Main St', 'New York', 'NY', '10001',  
'USA', '555-1234', '555-5678')
```

INSERT/UPSERT #1

```
CREATE (c:Customer {  
    customerID: 'ILYA1',  
    companyName: 'Acme Corp',  
    contactName: 'Ilya Verbitskiy',  
    contactTitle: 'Manager',  
    address: '123 Main St',  
    city: 'New York',  
    region: 'NY',  
    postalCode: '10001',  
    country: 'USA',  
    phone: '555-1234',  
    fax: '555-5678'  
});
```

INSERT/UPSERT #1

```
MERGE (c:Customer {customerID: 'ILYA1'})  
SET c.companyName = 'Acme Corp',  
    c.contactName = 'Ilya Verbitskiy',  
    c.contactTitle = 'Manager',  
    c.address = '123 Main St',  
    c.city = 'New York',  
    c.region = 'NY',  
    c.postalCode = '10001',  
    c.country = 'USA',  
    c.phone = '555-1234',  
    c.fax = '555-5678';
```

INSERT/UPSERT #2

```
BEGIN TRANSACTION
```

```
INSERT Orders(CustomerID, EmployeeID, OrderDate)  
VALUES ('ILYA1', 5, GETDATE())
```

```
DECLARE @LastOrderID INT = SCOPE_IDENTITY()  
DECLARE @ProductId INT  
SELECT @ProductId = ProductID FROM Products WHERE ProductName = 'Tofu'
```

```
INSERT OrderDetails(OrderID, ProductID, UnitPrice, Quantity, Discount)  
VALUES (@LastOrderID, @ProductId, 10, 8, 0)
```

```
COMMIT
```

INSERT/UPSERT #2

MATCH (c:Customer {customerID: 'ILYA1'})

MATCH (p:Product {productName: 'Tofu'})

MERGE (o:Order {orderID: "9090"}) SET o.orderDate = date()

MERGE u=(c)-[:PURCHASED]->(o)

MERGE v=(o)-[po:ORDERS {unitPrice: "10", quantity: 8, discount: "0", productID: p.productID, orderID: o.orderID}]->(p)

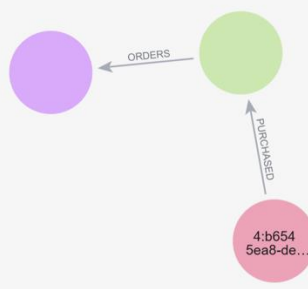
RETURN u, v;

INSERT/UPSERT #2

neo4j\$

```
1 MATCH (c:Customer {customerID: 'ILYAI'})
2 MATCH (p:Product {productName: 'Tofu'})
3 MERGE (o:Order {orderId: "9090"}) SET o.orderDate = date()
4 MERGE u=(c)-[:PURCHASED]->(o)
5 MERGE v=(o)-[:ORDERS {unitPrice: "10", quantity: 8, discount: "0", productID:
  p.productID, orderID: o.orderID}]->(p)
6 RETURN u, v
```

Graph Table RAW



Results overview

Nodes (3)

- * (3) Customer (1) Order (1) Product (1)

Relationships (2)

- * (2) ORDERS (1) PURCHASED (1)

Created 1 node, created 2 relationships, set 7 properties, added 1 label

Started streaming 1 record after 301 ms and completed after 357 ms.

UPDATE #1

Change Order.orderDate data type from String to Date

```
MERGE (o:Order)
```

```
SET o.orderDate = date(substring(o.orderDate, 0, 10));
```

UPDATE #2

Update the phone number.

```
UPDATE Customers SET Phone = '000-4321' WHERE CustomerID = 'ILYA1'
```

```
MERGE (c:Customer {customerID: 'ILYA1'})
```

```
SET c.phone = '000-4321';
```

UPDATE #3

Double the quantity of the order details record you inserted before

```
UPDATE od
SET Quantity = od.Quantity * 2
FROM OrderDetails od
INNER JOIN Orders o ON od.OrderID = o.OrderID
WHERE o.OrderID = 11084
```

```
MATCH (o:Order {orderID: "9090"})-[od:ORDERS]->(p:Product {productID: "14"})
SET od.quantity = od.quantity * 2
RETURN o, od, p;
```

DELETE #1

Delete the records you inserted before. Don't delete any other records!

DELETE #1

BEGIN TRANSACTION

DELETE od

FROM OrderDetails od

INNER JOIN dbo.Orders O on O.OrderID = od.OrderID

INNER JOIN Customers c ON o.CustomerID = c.CustomerID

WHERE c.CustomerID = 'ILYA1'

DELETE o

FROM Orders o INNER JOIN Customers c ON o.CustomerID = c.CustomerID

WHERE c.CustomerID = 'ILYA1'

DELETE Customers WHERE CustomerID = 'ILYA1'

COMMIT

DELETE #1

```
MATCH (c:Customer {customerID: "ILYA1"})-[:PURCHASED]->(o:Order)-[:ORDERS]->(p:Product)
```

```
DETACH DELETE c, o;
```

Let's keep in touch!

